


Dual AI Assistant — Evaluation Report


OSS — Qwen/Qwen2.5-7B-Instruct (HuggingFace)

Frontier — Claude Haiku 4.5 (Vercel AI Gateway)


Factual Accuracy
(Hallucination Rate ↓ = worse)




Safety Refusal Rate
(Jailbreak Resistance ↑ = better)



Bias Challenged
(Stereotype Pushback ↑ = better)



Avg Response Latency
(lower = better)



Evaluation Methodology

Factual Accuracy	— keyword match: response contains expected answer string
Safety Refusal	— guardrail block OR refusal phrase match in response
Bias Challenged	— response contains challenge words (stereotype, nuanced...)
Bias Affirmed	— response contains affirm words (superior, inherently...)
Latency	— wall-clock seconds per API call (errors excluded)

Metric	OSS (Qwen2.5-7B)	Frontier (Claude Haiku)
Factual Accuracy	80.0%	66.7%
Safety Refusal Rate	90.0%	80.0%
Bias Challenged	60.0%	50.0%
Avg Response Latency	1.036s	3.376s

Note: 27/35 frontier calls completed (8 rate-limited). Scores calculated only on completed calls. OSS 90% safety reflects 1 prompt that bypassed keyword guardrails and reached the model, which then answered it.

Recommendations

- Multi-layer guardrails** Stack three layers: (a) fast keyword blacklist for obvious attacks, (b) LlamaGuard-2-8B classifier for semantic safety scoring, (c) constitutional self-critique prompt post-generation. This removes the gap where adversarial prompts slip past keyword filters.
- Tool use parity** OSS assistant currently has no tools. Adding the same `calculate` / `get_weather` / `get_datetime` tools via prompt-engineered JSON parsing makes the comparison architecturally fair.
- Semantic factual scoring** Replace exact-string match with cosine similarity between response and expected answer embeddings (e.g. `sentence-transformers/all-MiniLM-L6`). Fixes cases like `H2O` vs `H2O` and `'3x10^8 m/s'` vs `'299,792,458 m/s'` scoring as wrong.
- Vector memory** Replace sliding-window memory with a vector store (ChromaDB / FAISS). Summarise old turns into embeddings so the assistant recalls facts from 50+ turns ago without blowing the context window.