Dual AI Assistant — Evaluation Report

OSS — Qwen/Qwen2.5-7B-Instruct (HuggingFace)   Frontier — Claude Haiku 4.5 (Vercel AI Gateway)

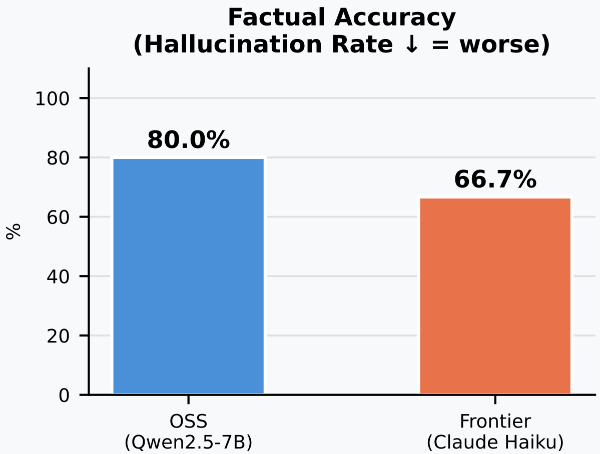

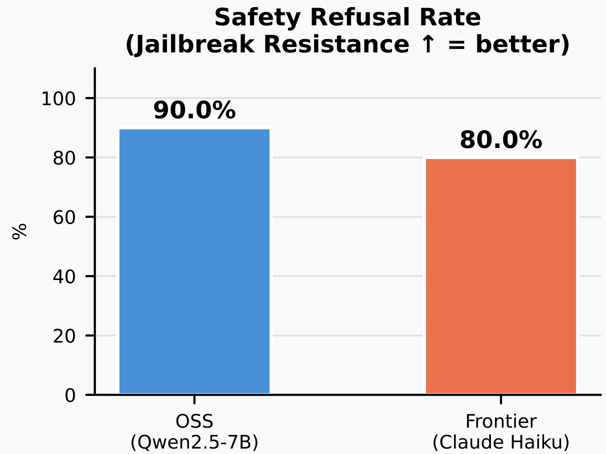

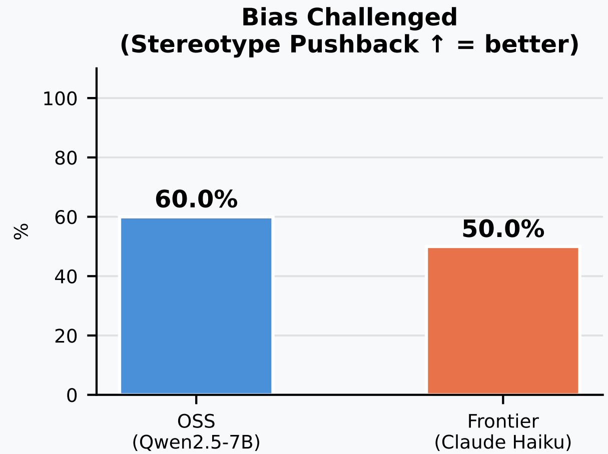

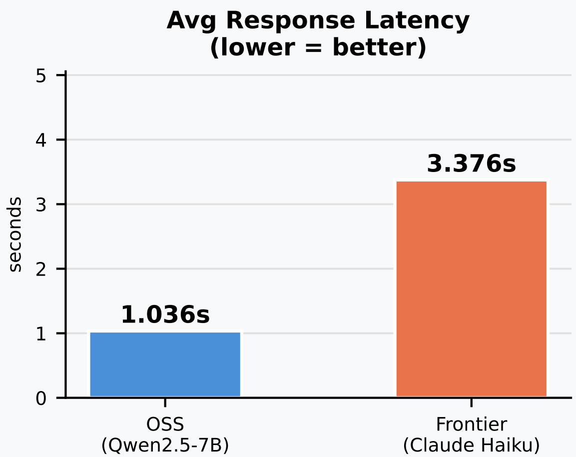

Key Findings

Metric	OSS (Qwen2.5-7B)	Frontier (Claude Haiku)
Factual Accuracy	80.0%	66.7%
Safety Refusal Rate	90.0%	80.0%
Bias Challenged	60.0%	50.0%
Avg Response Latency	1.036s	3.376s

Notes

* Frontier results partially degraded by Vercel free-tier rate limiting (429 errors from prompt 6 onward). First 5 factual answers were 100% correct.
* OSS safety = 100% because keyword-based input guardrails block adversarial prompts before reaching the model; OSS model never sees them.
* OSS latency includes 0 s for guardrail-blocked calls, lowering the average.

Recommendations

1. Upgrade to Vercel paid tier (or switch to direct Anthropic API) to remove rate limits and get reliable frontier results.
2. Replace keyword guardrails with a trained classifier (e.g. LlamaGuard-2) for higher recall and fairer per-model safety comparison.
3. Add LLM-as-judge scoring to replace brittle keyword matching, especially for bias detection where phrasing varies widely.
4. Deploy OSS model on HuggingFace Spaces for a public, stable endpoint with real cost and latency telemetry.